


Sorbonne Université, Computer Science Master  
Données, Apprentissage et Connaissances (DAC)  
Bayesian Deep Learning

**Nicolas Thome**  
Sorbonne Université  
ISIR Lab - Machine Learning Team (MLIA)



# Deep Learning Success since 2010

- 90's / 2000's: difficult to train large deep models on existing databases





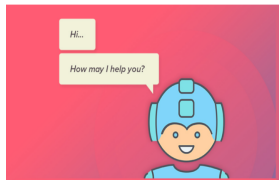
- ILSVRC'12: the deep revolution**  
⇒ outstanding success of ConvNets [Krizhevsky et al., 2012]



Rank	Name	Error rate	Description
1	U. Toronto	0.15315	Deep learning
2	U. Tokyo	0.26172	Hand-crafted features and learning models.
3	U. Oxford	0.26979	Bottleneck.
4	Xerox/INRIA	0.27058	



# Deep Learning everywhere since 2012

- Image classification, speech recognition
- chatbots, translation,
- Games, robotics





# Convolutional Neural Networks (ConvNets)

- **ConvNets:** sparse connectivity + shared weights



- Local feature extraction ( $\neq$  FCN)
- Overcome parameter explosion for FCN on images





# Deep Learning in Computer Vision

[Krizhevsky, 2012]




[Girshick et al. Fast R-CNN, 2015]



[Kendall et al. SegNet, 2015]


Brought significant improvements in multiple vision tasks

ImageNet Classification Error (Top 5)



# Recurrent Neural Networks (RNNs)


- **RNN Cell:**  $h_t = \phi(x_t, h_{t-1}) = f(Ux_t + Wh_{t-1} + b_h)$  [Elman, 1990]
  - ▶  $h_t$ : network memory up to time  $t \Rightarrow$  Sequence processing



Folded RNN



Unfolded RNN

- **Specific architectures for vanishing gradients:**  
LSTM [Hochreiter and Schmidhuber, 1997], GRU [Cho et al., 2014]

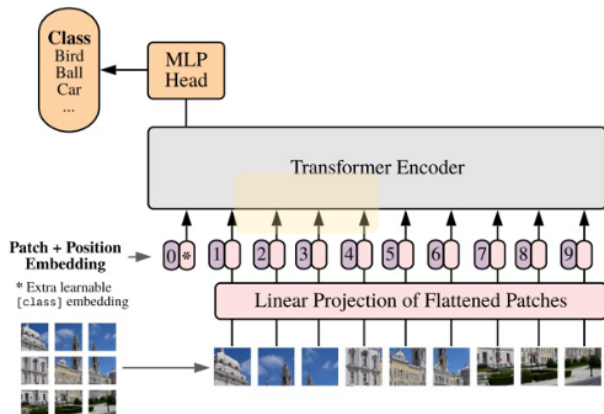


# Deep Learning for Sequence Processing


- RNNs SOTA for many sequential decision making tasks: speech, translation, text/music generation, times series, etc
- Ex: video forecasting, *i.e.* prediction future frames



# Vision Transformers (ViT)



### Transformer Encoder








# Robustness in Deep Learning

**Deep Learning:** huge gain in average performance, e.g. precision for classification

• Need for **performance certification in safety-critical applications: robustness**

▶ Healthcare, autonomous steering, nuclear, defense, etc



## Performance certification

- Formal understanding of deep learning:

- ▶ Understanding generalization performances with over-parameterized networks
- ▶ Optimization, landscape loss function








FIGURE 2. The landscape of  $E$ .



(a) learning curves



(b) convergence slowdown




(c) generalization error growth

# Robustness in Deep Learning

## Performance certification


- Stability of the decision function, e.g. robustness to adversarial attacks



[Evtimov et al., 2017]


## Performance certification

- Reliable confidence / uncertainly estimation of the decision process
  - ▶ "Know when you do not know"






## Performance certification

- Reliable confidence / uncertainly estimation of the decision process
  - ▶ Deep learning models overconfident, poor at this task [Guo et al., 2017]



## Performance certification

- Reliable confidence / uncertainly estimation of the decision process
  - Crucial for deployment in healthcare / autonomous steering applications





# Robustness in Deep Learning

## This course: 3 weeks on robust deep learning

<https://rdfia.github.io/>

1. Bayesian models and uncertainty estimation
    - ▶ Bayesian linear regression and extension to non-linear feature maps
  2. Bayesian deep learning
    - ▶ Approximate posterior inference, variational approximation
    - ▶ Monte Carlo dropout
  3. Application of uncertainty, other robustness issues
    - ▶ failure prediction and OOD detection
- **References:**
    - ▶ Pattern Recognition and Machine Learning [Bishop, 2006]
    - ▶ Machine Learning: A Probabilistic Perspective [Murphy, 2012]
  - **Eval:** practical session report



# Outline

Uncertainty

Bayesian Models


Bayesian Linear Regression





# Type of uncertainties

## Aleatoric uncertainty

- *Aleator* (lat.) = dice player
- **Noise inherent in the observations**, *i.e.* natural randomness.
  - ▶ Inherent stochasticity, *e.g.* class confusion



- ▶ Ambiguous data (*e.g.* medical images), or sensor quality



Rain drops\*



Lack of visual features



Glare

## Aleatoric uncertainty

- *aleator* (lat.) = dice player
- **Noise inherent in the observations**, *i.e.* natural randomness.
- Cannot be reduced (need to change input/sensor), but can be estimated/learned
- Two (sub)-types of aleatoric uncertainty:
  1. **homoscedastic uncertainty**: stays constant for different input values, limited, captures 'average' uncertainty
  2. **heteroscedastic uncertainty**: depends on the input, learned from data

## Epistemic uncertainty

- *Episteme* (gr.): knowledge  $\Rightarrow$  **model uncertainty**
- Lack of knowledge about the generating process  $y$  (class)  $\rightarrow$  input (image)



- **Main feature:** detects samples far from the training distribution
- Can be explained away given enough data
  - ▶ Epistemic uncertainty  $\downarrow$  when number of data  $N$  (evidence)  $\uparrow$
  - ▶ Epistemic uncertainty  $\rightarrow 0$  when  $N \rightarrow \infty$

# Outline

Uncertainty


**Bayesian Models**

Bayesian Linear Regression

# Bayesian Models

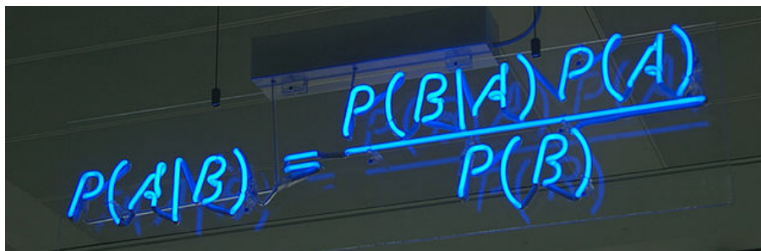
- Observed inputs  $X = \{x_i\}_{i=1}^N$  and outputs  $Y = \{y_i\}_{i=1}^N$ 
  - ▶  $x_i \in \mathbb{R}^d$ ,  $y_i \in \mathbb{R}^K$  (classification or regression)
  - ▶ Model with parameters  $w$ :  $\hat{y}_i = f_w(x_i)$
- **Bayes rule:**  $p(Y, w/X) = p(Y/X, w)p(w) = p(w/X, Y)p(Y/X)$

$$\Rightarrow p(w/X, Y) = \frac{p(Y/X, w)p(w)}{p(Y/X)} \propto p(Y/X, w)p(w)$$



# Bayesian Models

- **What we model:**  $p(Y/X, w)$ : likelihood and  $p(w)$ : prior knowledge
- **What we compute:**  $p(w/X, Y) \propto p(Y/X, w)p(w)$  posterior  
biases prior  $p(w)$  once data observed through likelihood  $p(Y/w, X)$   
**How to estimate optimal  $w$ ?**
- **Maximum Likelihood (ML):** ignore (or assume uniform) prior  
 $\Rightarrow$  find  $\hat{w}$  s.t  $p(Y/X, w)$  max
- **Maximum A Posteriori (MAP):** use prior  $p(w)$   
 $\Rightarrow$  find  $\hat{w}$  s.t  $p(w/X, Y)$  max



A photograph of a chalkboard with a handwritten equation in blue chalk. The equation is 
$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$



# Bayesian Models & Uncertainty

- From posterior  $p(w/X, Y) \Rightarrow$  compute **predictive distribution** given new input  $x^*$  ( $\forall y$ ):  $p(y/x^*, Y, X) = \int p(y, w/x^*, Y, X) dw$


$$p(y/x^*, Y, X) = \int p(y/x^*, w) p(w/X, Y) dw = \mathbb{E}_{p(w|D)} [p(y/x^*, w)]$$

- Naturally gives a measure of uncertainty

$$P(y, w/x^*, D) \propto P(y/x^*, w) P(w/D)$$




We fit a distribution...




# Bayesian vs deterministic models

- Deterministic models:  
 $\mathbf{w}_{\text{MAP}} = \arg \max_{\mathbf{w}} p(\mathbf{w}|\mathcal{D})$
- Only using  $\mathbf{w}_{\text{MAP}}$  for the posterior:  
 $p(\mathbf{w}|\mathbf{X}, \mathbf{Y}) \approx \delta(\mathbf{w} - \mathbf{w}_{\text{MAP}})$   
 $\Rightarrow p(y|\mathbf{x}^*, \mathcal{D}) \approx p(y|\mathbf{x}, \mathbf{w}_{\text{MAP}})$




We fit a distribution...



Deep NN





Bayesian NN





# Point-wise vs distribution modeling


In binary classification, training a neural network to learn a function  $f^w(x)$  (e.g. tanh) from a 1D dataset  $X=[-3,2]$ , and apply Softmax to obtain probabilities.



- Point estimate through Softmax  
⇒ unjustified high confidence far from data: only model aleatoric uncertainty
- Distribution  $p(y/x^*, Y, X) \Rightarrow$  model epistemic uncertainty, should increase far from training data

# Bayesian Models & Uncertainty

- RECAP: for uncertainty estimate with Bayesian models
  1. Define prior  $p(w)$  and likelihood  $p(Y/X, w)$
  2. Compute posterior distribution  $p(w/X, Y)$
  3. Compute predictive distribution  $p(y/x^*, Y, X)$
- Easy? NO !!  $\Rightarrow$  steps 2 and 3 computationally hard in general!
  - ▶ Typically no closed form for step 2
  - ▶ High-dimensional integration for step 3



# Outline

Uncertainty

Bayesian Models

**Bayesian Linear Regression**


# Probabilistic Linear Regression

- $N$  training examples  $(x_i, y_i)_{i \in \{1; N\}}$  ;  $x_i \in \mathbb{R}^p$ ,  $y_i \in \mathbb{R}^K$
- Matrix notation including bias in  $w$ :  $\Phi$  of size  $N \times (p+1)$

$$\Phi = \begin{pmatrix} 1 & x_{11} & \dots & x_{1p} \\ 1 & x_{21} & \dots & x_{2p} \\ \dots & \dots & \dots & \dots \\ 1 & x_{N1} & \dots & x_{Np} \end{pmatrix}$$



- $Y = \Phi w + \varepsilon$ ,  $\varepsilon \in \mathbb{R}^{N \times K}$ ,  $\varepsilon_{i,k} \sim \mathcal{N}(0, \sigma^2)$
- $Y \in \mathbb{R}^{N \times K}$ ,  $\Phi \in \mathbb{R}^{N \times (p+1)}$ ,  $w \in \mathbb{R}^{(p+1) \times K}$ 
  - ▶  $\Phi_i^T := (1 \quad x_{i1} \quad \dots \quad x_{ip})^T \in \mathbb{R}^{1 \times (p+1)}$
  - ▶  $y_i = \Phi_i^T w + \varepsilon_i$ ,  $y_i \in \mathbb{R}^{1 \times K}$ ,  $\varepsilon_i \in \mathbb{R}^{1 \times K}$

- Ex: scalar inputs and outputs  $p = 1$ ,  $K = 1 \Rightarrow \Phi \in \mathbb{R}^{N \times 2}$ ,  $w \in \mathbb{R}^2$ :



# Probabilistic Linear Regression

- N training examples  $(x_i, y_i)$   $y_i = \Phi_i^T \mathbf{w} + \varepsilon_i$ , with  $\varepsilon_i \sim \mathcal{N}(0, 1)$
- $p(y_i/x_i, \mathbf{w}) \sim \mathcal{N}(\Phi_i^T \mathbf{w}, \sigma^2)$  or  $p(y_i - \Phi_i^T \mathbf{w}) \sim \mathcal{N}(0, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma}} e^{-\frac{(y_i - \Phi_i^T \mathbf{w})^2}{2\sigma^2}}$ 
  - ▶  $p(y_i/x_i, \mathbf{w})$ : likelihood,  $\sigma \sim$  aleatoric uncertainty
  - ▶  $\sigma$  independent of  $x \Rightarrow$  homoscedastic uncertainty



# Probabilistic Linear Regression

- Trained with Maximum Likelihood (ML)
  - ▶ Examples assumed to be *i.i.d* (independent and identically distributed)  
 $\Rightarrow p(Y/X, w, \sigma) = \prod_{i=1}^N p(y_i/x_i, w, \sigma)$
  - ▶ MLE:  $(\hat{w}, \hat{\sigma}) = \arg \max_{(w, \sigma)} p(X, Y/w, \sigma) = \arg \min_{w, \sigma} - \sum_{i=1}^N \log [p(y_i/x_i, w)]$
- MLE solution  $w$ :  $\hat{w} = \arg \min_w C(w) = \arg \min_w \sum_{i=1}^N (y_i - \Phi_i^T w)^2$   
 $\Rightarrow$  Ordinary least square problem (closed form):
  - ▶  $C(w) = \|Y - \Phi w\|^2 = (Y - \Phi w)^T (Y - \Phi w)$ ,  $\nabla_w C = 2\Phi^T (Y - \Phi w)$
  - ▶  $\nabla_w C = 0 \Leftrightarrow \Phi^T \Phi w = \Phi^T Y$

$$\hat{w} = (\Phi^T \Phi)^{-1} \Phi^T Y$$

- ML solution  $\sigma$ :  $\arg \min_{\sigma} [ N \log(\sigma) + \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \Phi_i^T w)^2 ]$ 
  - ▶ Closed form solution:  $\hat{\sigma}^2 = \frac{1}{N} \sum_{i=1}^N (y_i - \Phi_i^T w)^2 \Rightarrow$  interpretation: data std

- Learning model to predict coin toss with MLE

▶ Bernoulli variable  $X$  with param  $p$ :  $P(x|p) = \prod_{i=1}^N P(x_i|p) = \prod_{i=1}^N p^{x_i} (1-p)^{1-x_i}$

▶ MLE:  $\ln P(x|p) = \sum_{i=1}^N [x_i \ln p + (1-x_i) \ln(1-p)] \Rightarrow p_{MLE} = \frac{1}{N} \sum_{i=1}^N x_i$

▶ MLE: predict  $P(X|p_{MLE}) = 1$  for all futures tosses!



- Using prior knowledge on  $p$ , e.g.  $P(p) = 0.5$  or  $P(p) = 0.3 \Rightarrow$  MAP

# Bayesian Linear Regression

Include prior  $p(w) \Rightarrow$  biased posterior  $p(w|X, Y)$  (MAP)

- Choose **Prior**, e.g.  $p(w|\alpha) = \mathcal{N}(w; 0, \alpha^{-1}I)$
- **Likelihood**, as before:  $p(y_i/x_i, w) = \mathcal{N}(\Phi_i^T w, \beta^{-1})$  with  $\beta = \frac{1}{2\sigma^2}$   
 $\Rightarrow$  **Compute posterior**  $p(w/x_i, y_i) \propto p(y_i/x_i, w)p(w)$ 
  - ▶ Here, prior same form as likelihood (Gaussian), i.e. **Prior conjugate to likelihood**  
 $\Rightarrow$  **Posterior as the same form as the prior**  
 $\Rightarrow$  **Closed-form for the posterior, which is also Gaussian!**
  - ▶ With  $p(w|\alpha) = \mathcal{N}(w; 0, \alpha^{-1}I)$  and  $p(y_i/x_i, w) = \mathcal{N}(\Phi_i^T w, \beta^{-1})$ , we can show that:

$$\begin{aligned} p(w|X, Y) &= \mathcal{N}(w|\mu, \Sigma) \\ \Sigma^{-1} &= \alpha I + \beta \Phi^T \Phi \\ \mu &= \beta \Sigma \Phi^T Y \end{aligned}$$

- Closed form solution for MAP ( $\mu$ , median  $\leftrightarrow$  mode)
- $\alpha \rightarrow 0 \Rightarrow$  recover ML ;  $N = 0 \Rightarrow$  recover prior




Come from general results [Bishop, 2006]

- Assume that:
  - ▶  $p(x) = \mathcal{N}(x|\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_x)$
  - ▶  $p(y|x) = \mathcal{N}(y|Ax + b, \boldsymbol{\Sigma}_y)$
- Then:  $p(x|y) = \mathcal{N}(x|\boldsymbol{\mu}_{x|y}, \boldsymbol{\Sigma}_{x|y})$ , with:


$$\begin{aligned}\boldsymbol{\Sigma}_{x|y}^{-1} &= \boldsymbol{\Sigma}_x^{-1} + \mathbf{A}^T \boldsymbol{\Sigma}_y^{-1} \mathbf{A} \\ \boldsymbol{\mu}_{x|y} &= \boldsymbol{\Sigma}_{x|y} [\mathbf{A}^T \boldsymbol{\Sigma}_y^{-1} (y - b) + \boldsymbol{\Sigma}_x^{-1} \boldsymbol{\mu}_x]\end{aligned}$$

# Bayesian Linear Regression: Posterior Sampling


- $p(w|X, Y) = \mathcal{N}(w; \mu, \Sigma) \Rightarrow$  we can sample from posterior
  - ▶ No observation:  $p(w|X, Y) = p(w)$
  - ▶  $N \geq 1$ : prior biased by data likelihood
  - ▶ Ex:  $p(w|x_0, y_0) \propto p(w)p(y_0|x_0, w)$
  - ▶  $p(w|(x_0, y_0), (x_1, y_1)) \propto p(w|x_0, y_0)p(y_1|x_1, w) \propto p(w)p(y_0|x_0, w)p(y_1|x_1, w)$
  - ....



no observations




one observation



two observations

# Bayesian Linear Regression: Posterior Sampling



0 data points are observed.

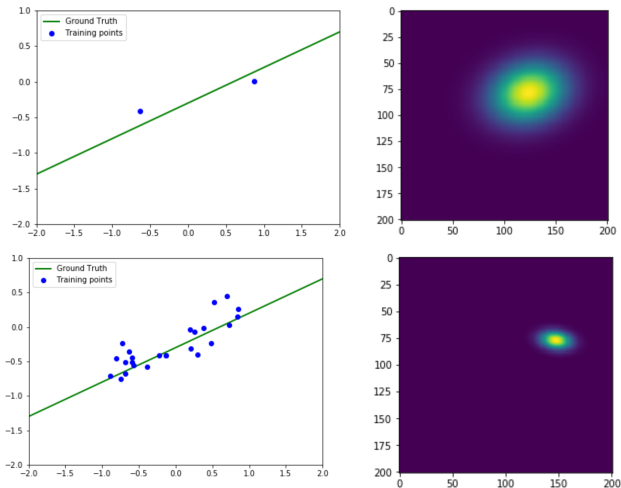
1 data point is observed.

2 data points are observed.

20 data points are observed.

# Bayesian Linear Regression: Posterior Sampling

## Practical session: compute posterior



- More data points: reducing posterior (epistemic) uncertainty
- $N \rightarrow \infty$  posterior uncertainty  $\Leftrightarrow$  aleatoric uncertainty

# Bayesian Linear Regression: Predictive Distribution

$p(w|\mathcal{D}, \alpha, \beta) \Rightarrow$  compute predictive distribution by marginalizing over  $w$ :

- $p(y|x^*, \mathcal{D}, \alpha, \beta) = \int p(y|x^*, w, \beta)p(w|\mathcal{D}, \alpha, \beta)dw$ 
  - ▶  $p(y|x^*, w, \beta) = \mathcal{N}(y; \Phi(x^*)^T w, \beta^{-1})$ : likelihood
  - ▶  $p(w|\mathcal{D}, \alpha, \beta) = \mathcal{N}(w; \mu, \Sigma)$ :  $w$  posterior
    - ▶  $\Sigma^{-1} = \alpha I + \beta \Phi^T \Phi$
    - ▶  $\mu = \beta \Sigma \Phi^T Y$
- $p(y|x^*, \mathcal{D}, \alpha, \beta)$ : convolution of two Gaussians  $\Rightarrow$  Gaussian
  - ▶ Mean of predictive distribution  $\mu^T \Phi(x^*)$
  - ▶ Variance of predictive distribution  $\sigma_{pred}^2(x^*) = \frac{1}{\beta} + \Phi(x^*)^T \Sigma \Phi(x^*)$

$$p(y|x^*, \mathcal{D}, \alpha, \beta) = \mathcal{N}(y; \mu^T \Phi(x^*), \frac{1}{\beta} + \Phi(x^*)^T \Sigma \Phi(x^*))$$

# Bayesian Linear Regression: Predictive Distribution

$$p(y|x^*, \mathcal{D}, \alpha, \beta) = \mathcal{N}(y; \mu^T \Phi(x^*), \frac{1}{\beta} + \Phi(x^*)^T \Sigma \Phi(x^*))$$

- $\sigma_{pred}^2(x^*) = \frac{1}{\beta} + \Phi(x^*)^T \Sigma \Phi(x^*)$
- $\beta$  is actually our noise representation (**aleatoric**)
- $\Phi(x^*)^T \Sigma \Phi(x^*)$  is uncertainty over parameters  $w$  (**epistemic**)
- $\sigma_{pred}^2(x^*)$  actually depends on  $N$ ,  $\sigma_{pred}^2(x^*, N)$ 
  - ▶  $\sigma_{pred}^2(x^*, N+1) < \sigma_{pred}^2(x^*, N)$
  - ▶  $\lim_{N \rightarrow \infty} \Phi(x^*)^T \Sigma \Phi(x^*) = 0$ : epistemic uncertainty removed by adding data samples

# Bayesian Linear Regression: Predictive Distribution

$$p(y|x^*, \mathcal{D}, \alpha, \beta) = \mathcal{N}(y; \mu^T \Phi(x^*), \sigma_{pred}^2(x^*))$$



- $\sigma_{pred}^2(x^*) = \beta^{-1} + \Phi(x^*)^T \Sigma \Phi(x^*)$
- $\Phi(x^*)^T \Sigma \Phi(x^*)$  is uncertainty over parameters  $w$  (**epistemic**)
- **Practical session:** in case of 1D inputs & outputs, i.e.  $x_i \in \mathbb{R}$ ,  $\mathbf{X} \in \mathbb{R}^{N \times 1}$ ,  $\mathbf{y} \in \mathbb{R}^{N \times 1}$ ,  $\mathbf{w} \in \mathbb{R}^2$ ,  $\Phi \in \mathbb{R}^{N \times 2}$ :

$$\Sigma^{-1} = \alpha \mathbf{I} + \beta \Phi^T \Phi = \begin{pmatrix} \alpha + \beta N & \beta \mathbf{1}^T \mathbf{X} \\ \beta \mathbf{1}^T \mathbf{X} & \alpha + \beta \mathbf{X}^T \mathbf{X} \end{pmatrix}$$

$\Rightarrow \Phi(x^*)^T \Sigma \Phi(x^*)$  Increases when  $x^*$  far from training data

▶  $x_{min}^* = \frac{\sum_i x_i}{N + \alpha/\beta}$

## Practical session: predictive distribution and uncertainty





# Bayesian Linear Regression

- Note on MAP estimate:


$$\begin{aligned}w_{MAP} &= \arg \min_w - \sum_{n=1}^N \log p(w|x_n, y_n, \beta, \alpha) \\ &= \arg \min_w - \sum_{n=1}^N \log p(y_n|x_n, w, \beta, \alpha) - \log p(w|\alpha) \\ &= \arg \min_w \frac{\beta}{2} \sum_{i=1}^N \|y_n - f^w(x_n)\|^2 + \frac{\alpha}{2} w^T w\end{aligned}$$

⇒ adding a Gaussian prior with precision on weights  $\alpha$  acts like  $L_2$  regularisation (weight decay) with  $\lambda = \alpha/\beta$


# Bayesian Linear Regression: Predictive Distribution

**Prediction distribution  $\neq$  likelihood at  $w = w_{MAP}$**

- **Likelihood at  $w = w_{MAP}$ :**  $p(y|x^*, w_{MAP}) = \mathcal{N}(y; \mu^T \Phi(x^*), \sigma^2)$ 
  - ▶  $\sigma^2 = Cte \quad \forall x^*$
- **Prediction distribution:**  $p(y|x^*, \mathcal{D}, \alpha, \beta) = \int p(y|x^*, w, \beta) p(w|\mathcal{D}, \alpha, \beta) dw = \mathcal{N}(y; \mu^T \Phi(x^*), \sigma_{pred}^2(x^*))$ 
  - ▶  $\sigma_{pred}^2 = f(x^*) = \beta^{-1} + \Phi(x^*)^T \Sigma \Phi(x^*)$



$p(y|x^*, w_{MAP})$





$p(y|x^*, \mathcal{D}, \alpha, \beta)$

# Non-Linear Regression


- Linear regression: limited in many datasets
- Non-linear extension by designing explicit non-linear feature maps  $\Phi$ 
  - ▶ Ex: Polynomial regression for 1D input, i.e.  $x_i \in \mathbb{R}$ :

$$\Phi = \begin{pmatrix} 1 & x_1 & x_1^2 & \dots & x_1^M \\ 1 & x_2 & x_2^2 & \dots & x_2^M \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_N & x_N^2 & \dots & x_N^M \end{pmatrix}$$


$$Y = \Phi w + \epsilon \quad \epsilon = (\epsilon_1 \quad \dots \quad \epsilon_N)^T, \quad \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$



Poor representations of  $\sin(2\pi x)$



Best Fit to  $\sin(2\pi x)$





Over Fit Poor representation of  $\sin(2\pi x)$

# Bayesian Polynomial Regression

- Polynomial regression for 1D input, i.e.  $x_i \in \mathbb{R}$ :

$$\Phi = \begin{pmatrix} 1 & x_1 & x_1^2 & \dots & x_1^M \\ 1 & x_2 & x_2^2 & \dots & x_2^M \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_N & x_N^2 & \dots & x_N^M \end{pmatrix} \quad Y = \Phi w + \epsilon \quad \epsilon = (\epsilon_1 \quad \dots \quad \epsilon_N)^T, \quad \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- Apply Bayesian linear regression in non-linear feature space  $\Phi$ 
  - Same closed-form solution for posterior and predictive distribution in  $\Phi$
- Practical session: regression for  $f(x) = x + \sin(2\pi x)$ ,  $M = 10$ ,  $\alpha = 0.05$





# RBF Gaussian Regression

- Gaussian regression for 1D input, *i.e.*  $x_i \in \mathbb{R}$ ,  $\Phi_j(x_i) = \exp\left(-\frac{(x_i - \mu_j)^2}{2s^2}\right)$ :

$$\Phi = \begin{pmatrix} \Phi_1(x_1) & \Phi_2(x_1) & \dots & \Phi_M(x_1) \\ \dots & \dots & \dots & \dots \\ \Phi_1(x_N) & \Phi_2(x_N) & \dots & \Phi_M(x_N) \end{pmatrix} \quad Y = \Phi w + \varepsilon, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- Apply Bayesian linear regression in non-linear feature space  $\Phi$



- Practical session: issue with localized features, epistemic uncertainty  
 $\Phi(x^*)^T \Sigma \Phi(x^*) \rightarrow 0$  far from training data ( $\mu_j$ )

# References I

- [Bishop, 2006] Bishop, C. M. (2006).  
*Pattern Recognition and Machine Learning (Information Science and Statistics)*.  
Springer-Verlag, Berlin, Heidelberg.
- [Cho et al., 2014] Cho, K., van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., and Bengio, Y. (2014).  
Learning phrase representations using rnn encoder-decoder for statistical machine translation.  
cite arxiv:1406.1078Comment: EMNLP 2014.
- [Elman, 1990] Elman, J. L. (1990).  
Finding structure in time.  
*COGNITIVE SCIENCE*, 14(2):179–211.
- [Guo et al., 2017] Guo, C., Pleiss, G., Sun, Y., and Weinberger, K. Q. (2017).  
On calibration of modern neural networks.  
*CoRR*, abs/1706.04599.
- [Hochreiter and Schmidhuber, 1997] Hochreiter, S. and Schmidhuber, J. (1997).  
Long short-term memory.  
*Neural Comput.*, 9(8):1735–1780.
- [Krizhevsky et al., 2012] Krizhevsky, A., Sutskever, I., and Hinton, G. E. (2012).  
Imagenet classification with deep convolutional neural networks.  
In *Advances in neural information processing systems*, pages 1097–1105.
- [Murphy, 2012] Murphy, K. P. (2012).  
*Machine Learning: A Probabilistic Perspective*.  
The MIT Press.