

Sorbonne Université, Computer Science Master
Données, Apprentissage et Connaissances (DAC)
Bayesian Deep Learning

Nicolas Thome
Sorbonne Université
ISIR Lab - Machine Learning Team (MLIA)



Beyond Bayesian Linear Regression

Bayesian Logistic Regression

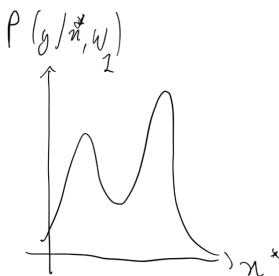
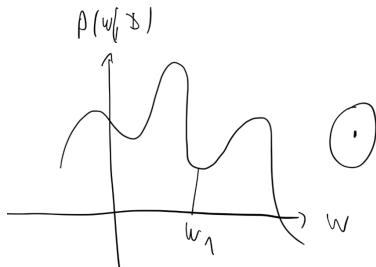
Bayesian Neural Networks

Monte Carlo Dropout

Beyond Bayesian Linear Regression

- Posterior distribution for parameters \mathbf{w} : $p(\mathbf{w}|\mathbf{X}, \mathbf{Y}) \propto p(\mathbf{Y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})$
- Predictive distribution $p(y|x^*, \mathcal{D}) = \int p(y|x^*, \mathbf{w})p(\mathbf{w}|\mathcal{D})d\mathbf{w}$, $(\mathbf{X}, \mathbf{Y}) := \mathcal{D}$

$$P(y, w | x^*, \mathcal{D}) \propto P(y | x^*, w) P(w | \mathcal{D})$$



$$P(y | x^*, \mathcal{D}) = \int P(w | \mathcal{D}) P(y | x^*, w) dw$$

Beyond Bayesian Linear Regression

- Posterior distribution for parameters \mathbf{w} : $p(\mathbf{w}|\mathbf{X}, \mathbf{Y}) \propto p(\mathbf{Y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})$
- Predictive distribution $p(y|\mathbf{x}^*, \mathcal{D}) = \int p(y|\mathbf{x}^*, \mathbf{w})p(\mathbf{w}|\mathcal{D})d\mathbf{w}$, $(\mathbf{X}, \mathbf{Y}) := \mathcal{D}$
- **Closed form for posterior $p(\mathbf{w}|\mathcal{D})$ and predictive distribution $p(y|\mathbf{x}^*, \mathcal{D})$: more the exception than the rule!**
- Slightly more complicated models : no closed form solution
 - ▶ Bayesian Logistic Regression
 - ▶ Simplest linear classification model
 - ▶ Likelihood not Gaussian
 - ▶ Neural network with one hidden layer in general
 - ▶ No closed form for regression and classification
 - ▶ And of course deep neural networks

No analytical expression for posterior $p(\mathbf{w}|\mathcal{D})$ and $p(y|\mathbf{x}^*, \mathcal{D})$ in general

⇒ **Approximation needed!**

- **Gaussian approximation for $p(\mathbf{w}|\mathcal{D})$**
 - ▶ **Ex: Laplace approximation** [MacKay, 1992]
 - ▶ Historically used for bayesian logistic regression
- **Monte Carlo methods:** sampling to directly evaluate integral $p(y|\mathbf{x}^*, \mathcal{D})$
 - ▶ **Metropolis-Hasting, Hamiltonian Monte Carlo** [Neal, 1996], **Expectation propagation** [Hernandez-Lobato and Adams, 2015, Jylänki et al., 2014]
- **Variational inference** [Hinton and van Camp, 1993, Graves, 2011, Blundell et al., 2015]: convert integration into optimization
 - ▶ Minimize KL divergence between $p(\mathbf{w}|\mathcal{D})$ and a proposed parametric function

Outline

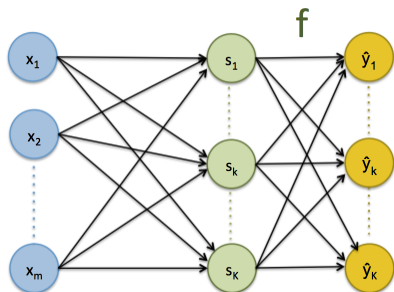
Beyond Bayesian Linear Regression

Bayesian Logistic Regression

Bayesian Neural Networks

Monte Carlo Dropout

Bayesian Logistic Regression (BLR)



- $\mathbf{s}_i = \mathbf{W}\mathbf{x}_i$
- Multi-class: $p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w}) = \hat{\mathbf{y}}_i$
 - ▶ $\hat{y}_{i,k} = \frac{\exp(s_k)}{\sum_k \exp(s_k)}$
- Binary case: $p(\mathbf{y}_i = 1 | \mathbf{x}_i, \mathbf{w}) = \sigma(s_i)$
 - ▶ σ sigmoid
 - ▶ $p(\mathbf{y}_i = -1 | \mathbf{x}_i, \mathbf{w}) = 1 - \sigma(s_i)$

$$p(\mathbf{w} | \mathbf{X}, \mathbf{Y}) \propto p(\mathbf{Y} | \mathbf{X}, \mathbf{w}) p(\mathbf{w})$$

- $p(\mathbf{Y} | \mathbf{X}, \mathbf{w}) = \prod_{i=1}^N p(\mathbf{y}_i = 1 | \mathbf{x}_i, \mathbf{w})$ not Gaussian anymore!

⇒ **no closed-form on posterior distribution $p(\mathbf{w} | \mathbf{X}, \mathbf{Y})$!**

Bayesian Logistic Regression training (MAP)

$$\begin{aligned}\mathbf{w}_{\text{MAP}} &= \arg \max_{\mathbf{w}} p(\mathbf{X}, \mathbf{Y} | \mathbf{w}) p(\mathbf{w}) = \arg \max_{\mathbf{w}} \prod_{n=1}^N p(y_n | \mathbf{x}_n, \mathbf{w}) p(\mathbf{w}) \\ &= \arg \min_{\mathbf{w}} \sum_{n=1}^N -\log(p(y_n | \mathbf{x}_n, \mathbf{w})) - \log(p(\mathbf{w}))\end{aligned}$$

- Gaussian priori: $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}; 0, \sigma_0^2 I)$;
- MAP BLR training with binary prediction:

$$\mathbf{w}_{\text{MAP}} = \arg \min_{\mathbf{w}} \sum_{n=1}^N \left(-y_n \log \sigma(\mathbf{w}^T \mathbf{x}_n + b) - (1 - y_n) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_n + b)) \right) + \frac{1}{2\sigma_0^2} \|\mathbf{w}\|_2^2$$

- Again: Gaussian prior \Leftrightarrow weight decay

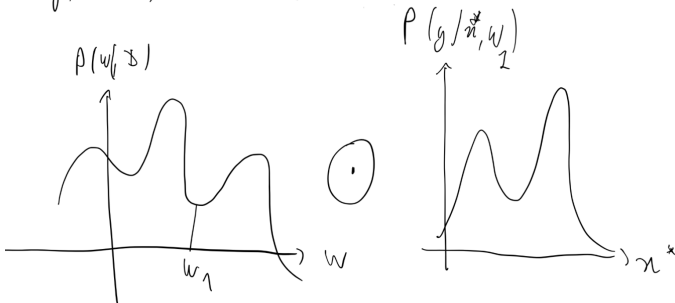
Bayesian Logistic Regression training (MAP)

$$\mathbf{w}_{\text{MAP}} = \arg \min_{\mathbf{w}} \sum_{n=1}^N \left(-y_n \log \sigma(\mathbf{w}^T \mathbf{x}_n + b) - (1 - y_n) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_n + b)) \right) + \frac{1}{2\sigma_0^2} \|\mathbf{w}\|_2^2$$

- \mathbf{w}_{MAP} with gradient descent
- Recap: we want to estimate predictive distribution:

$$p(y = 1 | \mathbf{x}^*, \mathcal{D}) = \int p(y = 1 | \mathbf{x}^*, \mathbf{w}) p(\mathbf{w} | \mathcal{D}) d\mathbf{w}$$

$$P(y, \mathbf{w} | \mathbf{x}^*, \mathcal{D}) \propto P(y | \mathbf{x}^*, \mathbf{w}) P(\mathbf{w} | \mathcal{D})$$



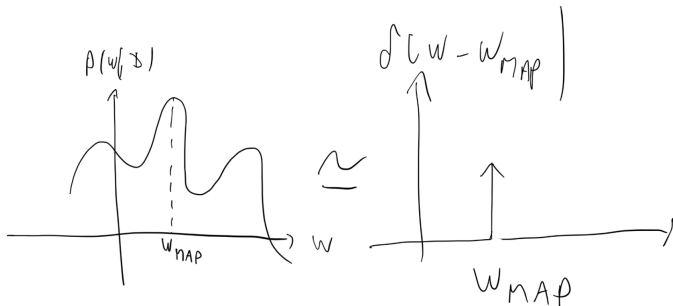
Bayesian Logistic Regression training (MAP)

$$\mathbf{w}_{\text{MAP}} = \arg \min_{\mathbf{w}} \sum_{n=1}^N \left(-y_n \log \sigma(\mathbf{w}^T \mathbf{x}_n + b) - (1 - y_n) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_n + b)) \right) + \frac{1}{2\sigma_0^2} \|\mathbf{w}\|_2^2$$

- \mathbf{w}_{MAP} with gradient descent
- Recap: we want to estimate predictive distribution:

$$p(y = 1 | \mathbf{x}^*, \mathcal{D}) = \int p(y = 1 | \mathbf{x}, \mathbf{w}) p(\mathbf{w} | \mathcal{D}) d\mathbf{w}$$

- ▶ Need full posterior distribution $p(\mathbf{w} | \mathbf{X}, \mathbf{Y})$, but posterior intractable
- ▶ $p(\mathbf{w} | \mathbf{X}, \mathbf{Y}) \approx \delta(\mathbf{w} - \mathbf{w}_{\text{MAP}}) \Rightarrow p(y = 1 | \mathbf{x}^*, \mathcal{D}) \approx p(y = 1 | \mathbf{x}, \mathbf{w}_{\text{MAP}})$



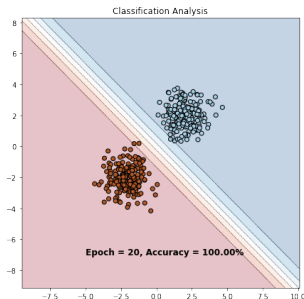
Bayesian Logistic Regression training (MAP)

$$\mathbf{w}_{\text{MAP}} = \arg \min_{\mathbf{w}} \sum_{n=1}^N \left(-y_n \log \sigma(\mathbf{w}^T \mathbf{x}_n + b) - (1 - y_n) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_n + b)) \right) + \frac{1}{2\sigma_0^2} \|\mathbf{w}\|_2^2$$

- \mathbf{w}_{MAP} with gradient descent
- Recap: we want to estimate predictive distribution:

$$p(y = 1 | \mathbf{x}^*, \mathcal{D}) = \int p(y = 1 | \mathbf{x}, \mathbf{w}) p(\mathbf{w} | \mathcal{D}) d\mathbf{w}$$

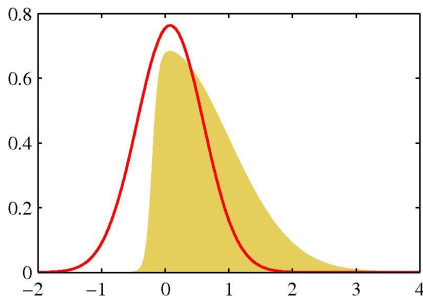
- ▶ Need full posterior distribution $p(\mathbf{w} | \mathbf{X}, \mathbf{Y})$, but posterior intractable
- ▶ $p(\mathbf{w} | \mathbf{X}, \mathbf{Y}) \approx \delta(\mathbf{w} - \mathbf{w}_{\text{MAP}}) \Rightarrow p(y = 1 | \mathbf{x}^*, \mathcal{D}) \approx p(y = 1 | \mathbf{x}, \mathbf{w}_{\text{MAP}})$



- $p(\mathbf{w} | \mathbf{X}, \mathbf{Y}) \approx \delta(\mathbf{w} - \mathbf{w}_{\text{MAP}})$: very coarse approximation:
- Uncertainty does not increase far from training data
- Need for more accurate approximations

Laplace Approximation for $p(\mathbf{w}|\mathbf{X}, \mathbf{Y})$

- Approximate $p(\mathbf{w}|\mathbf{X}, \mathbf{Y})$ by a normal distribution $q(\mathbf{w}) = \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$
- Fit the mean $\boldsymbol{\mu}$ of $q(\mathbf{w})$ to the mode of $p(\mathbf{w}|\mathbf{X}, \mathbf{Y})$
 - ▶ Mode of $p(\mathbf{w}) \Rightarrow \nabla_{\mathbf{w}} p(\mathbf{w}) = 0$
 - ▶ In practice, maximize log posterior (e.g. gradient ascent) \Rightarrow MAP: $\boldsymbol{\mu} = \mathbf{w}_{MAP}$
- Fit the inverse covariance $\boldsymbol{\Sigma}^{-1}$ of $q(\mathbf{w})$ to the Hessian of $p(\mathbf{w}|\mathbf{X}, \mathbf{Y})$ at $\boldsymbol{\mu} = \mathbf{w}_{MAP}$: $\boldsymbol{\Sigma}^{-1} = \nabla \nabla_{\mathbf{w}} p(\mathbf{w}|\mathbf{X}, \mathbf{Y})|_{\mathbf{w}=\mathbf{w}_{MAP}}$



from [Bishop, 2006]

- Laplace limitation: approximation at a single value of $p(\mathbf{w}|\mathbf{X}, \mathbf{Y})$, ignores global properties

Predictive Distribution $p(y|\mathbf{x}^*, \mathcal{D})$ for BLR

RECAP, with $\mathcal{D} = \mathbf{X}, \mathbf{Y}$:

$$p(y|\mathbf{x}^*, \mathcal{D}) = \int p(y|\mathbf{x}^*, \mathbf{w})p(\mathbf{w}|\mathcal{D})d\mathbf{w}$$

- Posterior approximation by normal $p(\mathbf{w}|\mathcal{D}) \approx q(\mathbf{w}) = \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$:

$$p(y|\mathbf{x}^*, \mathcal{D}) \approx \int p(y|\mathbf{x}^*, \mathbf{w})q(\mathbf{w})d\mathbf{w}$$

- However, likelihood $p(y|\mathbf{x}^*, \mathcal{D})$ still not Gaussian

⇒ **Intractable posterior distribution** $p(y|\mathbf{x}^*, \mathcal{D})!$

Predictive Distribution $p(y|x^*, \mathcal{D})$ for BLR

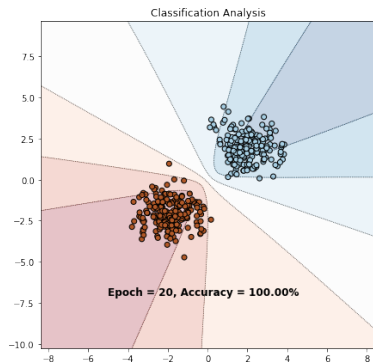
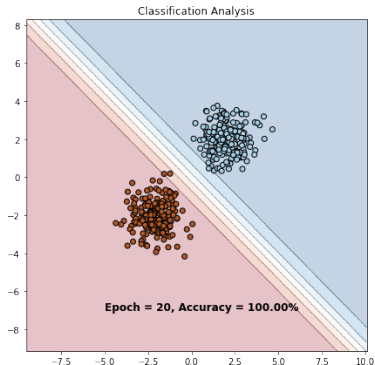
- **Option 1:** use Monte Carlo (MC) sampling

- ▶ Binary case (simple x-class extension): $p(y = 1|x^*, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}^*)$, σ sigmoid

$$p(y = 1|x^*, \mathcal{D}) \approx \sum_{s=1}^S \sigma\left((\mathbf{w}^s)^T \mathbf{x}^*\right) \quad \mathbf{w}^s \sim q(\mathbf{w}) = \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$$

- ▶ Easy to sample from Gaussian $q(\mathbf{w})$

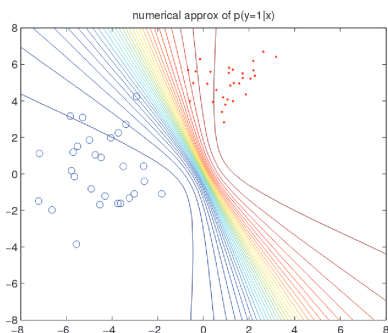
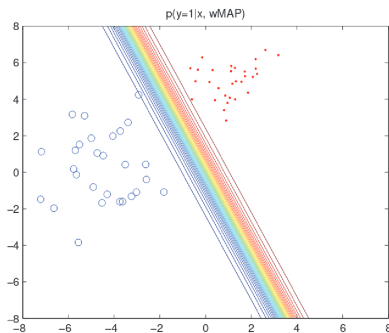
- **Practical session:** MAP solution for LR vs BLR (Laplace & MC sampling)



Predictive Distribution $p(y|\mathbf{x}^*, \mathcal{D})$ for BLR

- $p(y|\mathbf{x}^*, \mathcal{D}) \approx \int p(y|\mathbf{x}^*, \mathbf{w})q(\mathbf{w})d\mathbf{w}$ intractable
- **Option 2:** (binary case): $p(y|\mathbf{x}^*, \mathcal{D}) \approx \int \sigma(\mathbf{w}^T \mathbf{x}^*)q(\mathbf{w})d\mathbf{w}$; $q(\mathbf{w}) = \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$
- Convolution of sigmoid with Gaussian still intractable
 - ▶ Approximate $\sigma(\mathbf{w}^T \mathbf{x}^*)$ by probit: $\sigma(a) \approx \Phi(\lambda a)$, $\lambda^2 = \pi/8$
 - ▶ Convolution of probit with Gaussian \Rightarrow probit:

$$p(y|\mathbf{x}^*, \mathcal{D}) \approx \int \Phi(\lambda \mathbf{w}^T \mathbf{x}^*) \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \Phi\left(\frac{\mu_f}{\sqrt{\lambda^{-2} + \sigma_f^2}}\right) \quad \mu_f = \boldsymbol{\mu}^T \mathbf{x}^* \quad \sigma_f^2 = \mathbf{x}^{*T} \boldsymbol{\Sigma} \mathbf{x}^*$$



from [Murphy, 2012]

Outline

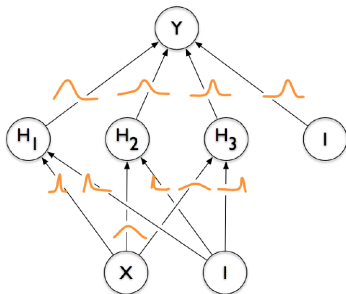
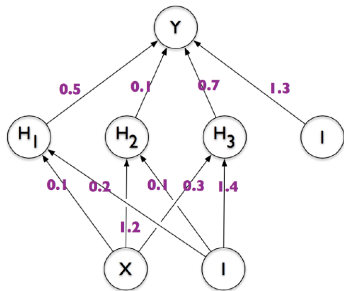
Beyond Bayesian Linear Regression

Bayesian Logistic Regression

Bayesian Neural Networks

Monte Carlo Dropout

Bayesian Neural Networks (BNN)



Credit: [Blundell et al., 2015]

- Standard NN: $\mathbf{y}_i = f^{\mathbf{w}}(\mathbf{x}_i)$, **Bayesian NN**: $p(\mathbf{y}_i | \mathbf{x}_i, \mathcal{D})$
 - ▶ Define prior over weights, e.g. $p(\mathbf{w}) = \mathcal{N}(\mathbf{w} | 0, \alpha^{-1} \mathcal{I})$ (point estimate for bias)
 - ▶ In practice, typically separate variance $\sigma^2 = \alpha^{-1}$ for each layer
 - ▶ Define likelihood, $p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$, e.g. for regression $p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w}) = \mathcal{N}(\mathbf{y}_i; f^{\mathbf{w}}(\mathbf{x}_i), \beta^{-1})$
 - ▶ **Goal**: compute posterior $p(\mathbf{w} | \mathbf{X}, \mathbf{Y}) = \prod_{i=1}^N p(\mathbf{w} | \mathbf{x}_i, \mathbf{y}_i, \beta) \propto p(\mathbf{w}) \prod_{i=1}^N p(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$

Bayesian Neural Networks (BNN)

$$p(\mathbf{w}|\mathbf{X}, \mathbf{Y}) \propto p(\mathbf{w}) \prod_{i=1}^N p(\mathbf{y}_i|\mathbf{x}_i, \mathbf{w})$$

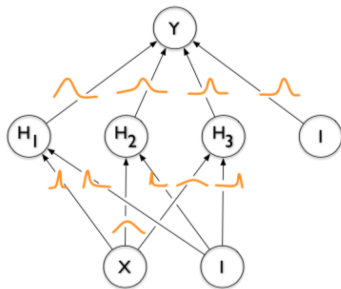
- **With Bayesian Neural networks, even with:**

- ▶ Gaussian prior $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|0, \alpha^{-1}\mathcal{I})$
- ▶ Gaussian likelihood, e.g. regression $p(\mathbf{y}_i|\mathbf{x}_i, \mathbf{w}) = \mathcal{N}(\mathbf{y}_i; f^{\mathbf{w}}(\mathbf{x}_i), \beta^{-1})$
- ▶ Posterior $p(\mathbf{w}|\mathbf{X}, \mathbf{Y}, \beta) \propto p(\mathbf{w}) \prod_{i=1}^N p(\mathbf{y}_i|\mathbf{x}_i, \mathbf{w})$ is **NOT Gaussian!!**

- **Non-linear dependence of $f^{\mathbf{w}}(\mathbf{x})$ on \mathbf{w} !**

- RECAP:

- ▶ $p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_x)$
- ▶ $p(\mathbf{y}|\mathbf{x}) = \mathcal{N}(\mathbf{y}|\mathbf{Ax} + \mathbf{b}, \boldsymbol{\Sigma}_y)$
 - ▶ Linear dependence $\mathbf{Ax} + \mathbf{b}$ required
- ▶ Then: $p(\mathbf{x}|\mathbf{y}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_{x|\mathbf{y}}, \boldsymbol{\Sigma}_{x|\mathbf{y}})$
 - ▶ Not true for BNNs!



Credit: [Blundell et al., 2015]

Posterior Inference: MCMC Carlo Sampling

The true predictive distribution $p(y|\mathbf{x}^*, \mathcal{D})$ cannot be evaluated analytically

$$p(y|\mathbf{x}^*, \mathcal{D}) = \int p(y|\mathbf{x}^*, \mathbf{w})p(\mathbf{w}|\mathcal{D})d\mathbf{w}$$

- Monte Carlo estimation of the integral:

$$p(y|\mathbf{x}^*, \mathcal{D}) \approx \frac{1}{S} \sum_{s=1}^S p(y|\mathbf{x}^*, \mathbf{w}^s) \quad \mathbf{w}^s \sim p(\mathbf{w}|\mathcal{D})$$

- Can't sample exactly from $p(\mathbf{w}|\mathcal{D})$, **BUT approximate sampling using Markov chain Monte Carlo (MCMC) possible !**
 - ▶ Metropolis-Hasting (MH), Hamiltonian Monte Carlo (HMC) [Neal, 1996]
- **Works well, accurate posterior inference in BNNs**
- **Main drawback: does not scale to large datasets**
 - ▶ Computing likelihood for MH/HMC acceptance step requires the whole dataset

Variational Inference (VI)

The true posterior $p(\mathbf{w}|\mathbf{X}, \mathbf{Y})$ cannot usually be evaluated analytically

- Defining an **approximating variational distribution** $q_{\theta}(\mathbf{w})$, parameterized by θ
- **Minimizing its KL divergence with the true posterior:**

$$KL(q_{\theta}(\mathbf{w})||p(\mathbf{w}|\mathbf{X}, \mathbf{Y})) = \int q_{\theta}(\mathbf{w}) \log \frac{q_{\theta}(\mathbf{w})}{p(\mathbf{w}|\mathbf{X}, \mathbf{Y})} d\mathbf{w}$$

- **Computing approximate predictive distribution:** $p(\mathbf{w}|\mathbf{X}, \mathbf{Y}) \leftarrow q_{\theta^*}(\mathbf{w})$:

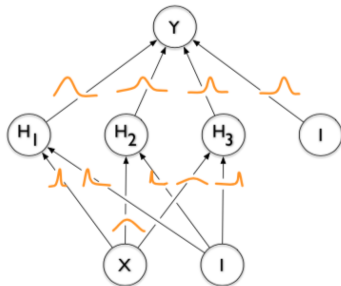
$$p(y|x^*, \mathbf{X}, \mathbf{Y}) \approx \int p(y|x^*, \mathbf{w})q_{\theta^*}(\mathbf{w})d\mathbf{w}$$

Variational Inference (VI)

- Recap: BNN prior, e.g. Gaussian: $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|0, \alpha^{-1}\mathcal{I})$
- **Variational approximate posterior** $q_{\theta}(\mathbf{w})$, e.g. fully factorized Gaussian:

$$q_{\theta}(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\theta) = \mathcal{N}(\mathbf{w}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \prod_{i=1}^D \mathcal{N}(w_i|\mu_i, \sigma_i)$$

- Each weight of the network w_j has its own mean μ_j and variance σ_j
 - ▶ $\theta = \{(\mu_j, \sigma_j)\}_{j \in \{1; D\}}$: **variational parameters**



Credit: [Blundell et al., 2015]

Variational Inference (VI): ELBO

$$\begin{aligned}KL(q_{\theta}(\mathbf{w})\|p(\mathbf{w}|\mathbf{X}, \mathbf{Y})) &= \int q_{\theta}(\mathbf{w}) \log \frac{q_{\theta}(\mathbf{w})}{p(\mathbf{w}|\mathbf{X}, \mathbf{Y})} d\mathbf{w} = - \int q_{\theta}(\mathbf{w}) \log \frac{p(\mathbf{w}|\mathbf{X}, \mathbf{Y})}{q_{\theta}(\mathbf{w})} d\mathbf{w} \\ &= - \int q_{\theta}(\mathbf{w}) \log \frac{p(\mathbf{Y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})}{q_{\theta}(\mathbf{w})p(\mathbf{Y}|\mathbf{X})} d\mathbf{w} \\ &= - \int q_{\theta}(\mathbf{w}) \log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) d\mathbf{w} + \int q_{\theta}(\mathbf{w}) \log \frac{q_{\theta}(\mathbf{w})}{p(\mathbf{w})} + \log p(\mathbf{Y}|\mathbf{X}) \\ &= - \int q_{\theta}(\mathbf{w}) \log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) d\mathbf{w} + KL(q_{\theta}(\mathbf{w})\|p(\mathbf{w})) + \log p(\mathbf{Y}|\mathbf{X})\end{aligned}$$

- $\Rightarrow \boxed{KL(q_{\theta}(\mathbf{w})\|p(\mathbf{w}|\mathbf{X}, \mathbf{Y})) = -\mathcal{L}_{VI}(\mathbf{X}, \mathbf{Y}, \theta) + \log p(\mathbf{Y}|\mathbf{X})}$

- ▶ $\mathcal{L}_{VI}(\mathbf{X}, \mathbf{Y}, \theta)$: Evidence Lower Bound (ELBO)

$$\mathcal{L}_{VI}(\theta) = \int q_{\theta}(\mathbf{w}) \log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) d\mathbf{w} - KL(q_{\theta}(\mathbf{w})\|p(\mathbf{w}))$$

- ▶ $\mathcal{L}_{VI}(\theta) = \log p(\mathbf{Y}|\mathbf{X}) - KL(q_{\theta}(\mathbf{w})\|p(\mathbf{w}|\mathbf{X}, \mathbf{Y})) \leq \log p(\mathbf{Y}|\mathbf{X})$

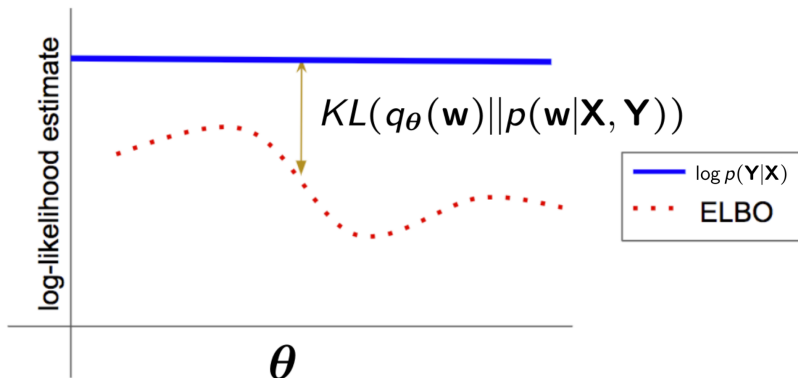
$$p(\mathbf{w}|\mathbf{X}, \mathbf{Y}) = \frac{p(\mathbf{Y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})}{p(\mathbf{Y}|\mathbf{X})}$$

ELBO Illustration

- $\mathcal{L}_{VI}(\mathbf{X}, \mathbf{Y}, \theta)$: Evidence Lower Bound (ELBO)

$$\mathcal{L}_{VI}(\theta) = \int q_{\theta}(\mathbf{w}) \log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) d\mathbf{w} - KL(q_{\theta}(\mathbf{w})||p(\mathbf{w}))$$

- $\mathcal{L}_{VI}(\theta) = \log p(\mathbf{Y}|\mathbf{X}) - KL(q_{\theta}(\mathbf{w})||p(\mathbf{w}|\mathbf{X}, \mathbf{Y})) \leq \log p(\mathbf{Y}|\mathbf{X})$

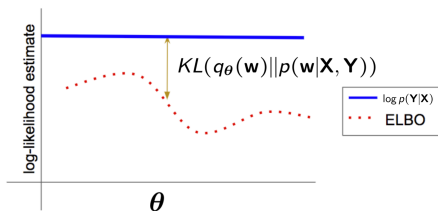


Variational Inference (VI): ELBO

- $\mathcal{L}_{VI}(\theta) = \log p(\mathbf{Y}|\mathbf{X}) - KL(q_{\theta}(\mathbf{w})||p(\mathbf{w}|\mathbf{X}, \mathbf{Y}))$
- \Rightarrow Minimizing $KL(q_{\theta}(\mathbf{w})||p(\mathbf{w}|\mathbf{X}, \mathbf{Y})) \Leftrightarrow$ maximizing $\mathcal{L}_{VI}(\theta)$ w.r.t $q_{\theta}(\mathbf{w})$:

$$\begin{aligned}\mathcal{L}_{VI}(\theta) &= \int q_{\theta}(\mathbf{w}) \log p(\mathbf{Y}|\mathbf{X}, \mathbf{w}) d\mathbf{w} - KL(q_{\theta}(\mathbf{w})||p(\mathbf{w})) \leq \log p(\mathbf{Y}|\mathbf{X}) \\ &= \mathbb{E}_{q_{\theta}(\mathbf{w})}[\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})] - KL(q_{\theta}(\mathbf{w})||p(\mathbf{w})) \\ &= \sum_{i=1}^N \int q_{\theta}(\mathbf{w}) \log p(\mathbf{y}_i|f^{\mathbf{w}}(\mathbf{x}_i)) d\mathbf{w} - KL(q_{\theta}(\mathbf{w})||p(\mathbf{w}))\end{aligned}$$

- **Exp. log likelihood** $\mathbb{E}_{q_{\theta}(\mathbf{w})}[\log p(\mathbf{Y}|\mathbf{X}, \mathbf{w})]$: max $\Leftrightarrow q_{\theta}(\mathbf{w})$ explain data well
- **Prior KL** $KL(q_{\theta}(\mathbf{w})||p(\mathbf{w}))$: min $\Leftrightarrow q_{\theta}(\mathbf{w})$ as close as possible to $p(\mathbf{w})$ prior



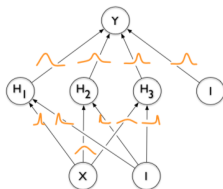
Variational Inference: Training

- Variational Bayesian NN training: computing derivatives of \mathcal{L}_{VI} w.r.t variational parameters θ

$$\mathcal{L}_{VI}(\theta) = \sum_{i=1}^N \int q_{\theta}(\mathbf{w}) \log p(\mathbf{y}_i | f^{\mathbf{w}}(\mathbf{x}_i)) d\mathbf{w} - KL(q_{\theta}(\mathbf{w}) || p(\mathbf{w}))$$

- RECAP: approximate variational posterior, e.g. fully factorize Gaussian:

$$q_{\theta}(\mathbf{w}) = \prod_{i=1}^D \mathcal{N}(w_j | \mu_j, \sigma_j)$$



- **Prior KL** $KL(q_{\theta}(\mathbf{w}) || p(\mathbf{w}))$: often can be integrated analytically, e.g. with Gaussian functions for prior $p(\mathbf{w})$ and posterior approximation $q_{\theta}(\mathbf{w})$
- **Expected log likelihood** $\mathbb{E}_{q_{\theta}(\mathbf{w})}[\log p(\mathbf{Y} | \mathbf{X}, \mathbf{w})]$: no close-form solution in general, requires tractable calculations over the entire dataset
⇒ estimation by sampling

Stochastic Variational Inference (VI): Training

$$\mathcal{L}_{VI}(\theta) = \sum_{i=1}^N \int q_{\theta}(\mathbf{w}) \log p(\mathbf{y}_i | f^{\mathbf{w}}(\mathbf{x}_i)) d\mathbf{w} - KL(q_{\theta}(\mathbf{w}) || p(\mathbf{w}))$$

- **Scalable gradient computation: batch sampling** [Graves, 2011]
 - ▶ \mathcal{L}_{VI} linearly decomposes into training examples, unbiased gradient estimator
- **Modern solutions: approximate integral with MC integration** $\hat{\mathbf{w}}_i \sim q_{\theta}(\mathbf{w})$
 - ▶ Sample $\log p(\mathbf{y}_i | f^{\hat{\mathbf{w}}_i}(\mathbf{x}_i))$, $\hat{\mathbf{w}}_i \sim q_{\theta}(\mathbf{w})$

$$\mathcal{L}_{VI}(\theta) = \sum_{i \in S} \log p(\mathbf{y}_i | f^{\hat{\mathbf{w}}_i}(\mathbf{x}_i)) - KL(q_{\theta}(\mathbf{w}) || p(\mathbf{w}))$$

- **Issue: computing gradient wrt variational parameters** $\frac{\partial}{\partial \theta} \log p(\mathbf{y}_i | f^{\hat{\mathbf{w}}_i}(\mathbf{x}_i))$
- **Problem: sampling** $\hat{\mathbf{w}}_i \sim q_{\theta}(\mathbf{w})$ **depends on variational parameters** θ

- ▶ **Solution, easy cases: re-parametrization** $\mathbf{w} = g(\theta, \epsilon)$
 - ▶ Where g deterministic and ϵ independent of θ - As in VAE [Kingma and Welling, 2014]
 - ▶ **Crucial point: sampling fully in ϵ , independent of θ**
- ▶ **Gaussian ex:** $\theta = \{\theta_j\}$, $\theta_j = (\mu_j, \sigma_j)$: $\hat{\mathbf{w}}_j \sim \mathcal{N}(\mathbf{w}_j | \mu_j, \sigma_j)$
 - ▶ $\mathbf{w}_j = g((\mu_j, \sigma_j), \epsilon_j) = \mu_j + \sigma_j \epsilon_j$: $\hat{\mathbf{w}}_j \sim \mathcal{N}(\mathbf{w}_j | \mu_j, \sigma_j) \Leftrightarrow \hat{\epsilon}_j \sim \mathcal{N}(\epsilon_j | 0, 1)$

$$\mathcal{L}_{VI}(\theta) = \sum_{i \in S} \log p(\mathbf{y}_i | f^{g(\theta, \hat{\epsilon}_i)}(\mathbf{x}_i)) - KL(q_{\theta}(\mathbf{w}) || p(\mathbf{w}))$$

Stochastic Variational Inference (VI): Training

Algorithm 1 Minimise divergence between $q_\theta(\boldsymbol{\omega})$ and $p(\boldsymbol{\omega}|X, Y)$

- 1: Given dataset \mathbf{X}, \mathbf{Y} ,
- 2: Define learning rate schedule η ,
- 3: Initialise parameters θ randomly.
- 4: **repeat**
- 5: Sample M random variables $\hat{\boldsymbol{\epsilon}}_i \sim p(\boldsymbol{\epsilon})$, S a random subset of $\{1, \dots, N\}$ of size M .
- 6: Calculate stochastic derivative estimator w.r.t. θ :

$$\widehat{\Delta\theta} \leftarrow -\frac{N}{M} \sum_{i \in S} \frac{\partial}{\partial \theta} \log p(\mathbf{y}_i | \mathbf{f}^{g(\theta, \hat{\boldsymbol{\epsilon}}_i)}(\mathbf{x}_i)) + \frac{\partial}{\partial \theta} \text{KL}(q_\theta(\boldsymbol{\omega}) || p(\boldsymbol{\omega})).$$

- 7: Update θ :
 $\theta \leftarrow \theta + \eta \widehat{\Delta\theta}$.
 - 8: **until** θ has converged.
-

Bayesian Neural Networks: Predictive Distribution

- Gaussian approximation of posterior $q_{\theta}(\mathbf{w}) \approx p(\mathbf{w}|\mathbf{X}, \mathbf{Y})$ e.g. VI or Laplace
- Predictive distribution: $p(\mathbf{y}|\mathbf{x}^*, \mathcal{D}) \approx \int p(\mathbf{y}|\mathbf{x}^*, \mathbf{w})q_{\theta}(\mathbf{w})d\mathbf{w}$
- **Even with $q_{\theta}(\mathbf{w})$ Gaussian, no closed-form for $p(\mathbf{y}|\mathbf{x}^*, \mathcal{D})!$**
 - ▶ Again due to non-linear dependence of $\mathbf{y}(\mathbf{x}^*, \mathbf{w})$ wrt \mathbf{w}

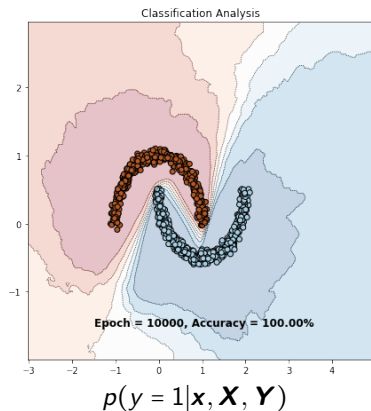
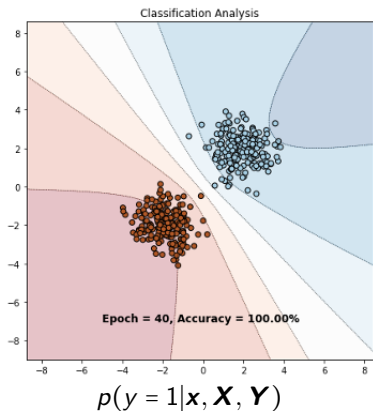
Solutions:

1. MC sampling, easy to sample from $q_{\theta}(\mathbf{w})$
2. Perform Taylor expansion of $\mathbf{y}(\mathbf{x}^*, \mathbf{w})$ around \mathbf{w}_{MAP} for regression^a:
$$y(\mathbf{x}, \mathbf{w}) \approx y(\mathbf{x}, \mathbf{w}_{MAP}) + \frac{\partial y}{\partial \mathbf{w}} \Big|_{\mathbf{w}=\mathbf{w}_{MAP}} (\mathbf{w} - \mathbf{w}_{MAP})$$
 - $p(\mathbf{y}|\mathbf{x}^*, \mathbf{w}) \approx \mathcal{N}(y|\mu_y; \beta^{-1})$, $\mu_y = y(\mathbf{x}, \mathbf{w}_{MAP}) + \frac{\partial y}{\partial \mathbf{w}} \Big|_{\mathbf{w}=\mathbf{w}_{MAP}} (\mathbf{w} - \mathbf{w}_{MAP})$
 - Closed form solution for $p(\mathbf{y}|\mathbf{x}^*, \mathcal{D})$
 - ▶ $p(\mathbf{y}|\mathbf{x}^*, \mathcal{D}) = \mathcal{N}(y|y(\mathbf{x}, \mathbf{w}_{MAP}); \sigma^2(x))$
 - ▶ $\sigma^2(x) = \beta^{-1} + \mathbf{g}^T \mathbf{A}^{-1} \mathbf{g}$, \mathbf{g} gradient and \mathbf{A} Hessian at \mathbf{w}_{MAP}

^aFor classification, logit Taylor expansion, see 5.7.1 in [Bishop, 2006]

Practical session

- Own implementation of VI for:
 - ▶ Bayesian logistic regression
 - ▶ Neural network for non-linear classification (moons)



Outline

Beyond Bayesian Linear Regression

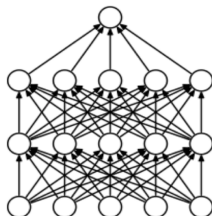
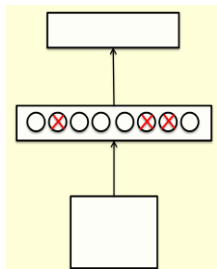
Bayesian Logistic Regression

Bayesian Neural Networks

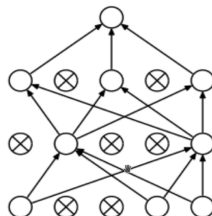
Monte Carlo Dropout

Dropout [Hinton et al., 2012]

- Randomly omit each hidden unit with probability p , e.g. $p = 0.5$
- **Regularization technique**, limits over-fitting (better generalization)
 - ▶ Prevent co-adaptation
 - ▶ May be viewed as averaging over many NN
 - ▶ Slower convergence



Standard Neural Net

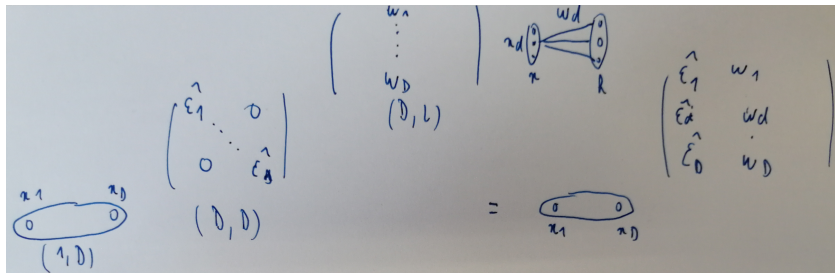


After applying dropout.

Credits: Geoffrey E. Hinton, NIPS 2012

Dropout as a variational inference [Gal, 2016]

- Input $\mathbf{x} \in \mathbb{R}^{D^a}$, latent vector $\mathbf{h} \in \mathbb{R}^L$
 - ▶ First layer: $\mathbf{h} = \sigma(\mathbf{x}\mathbf{W}_1)$, σ non-linearity
- **Dropout sampling:** in input : $\mathbf{x} \odot \hat{\boldsymbol{\varepsilon}}$
 - ▶ $\hat{\boldsymbol{\varepsilon}} = \{\hat{\varepsilon}_i^1\}_{i \in \{1;D\}}$ $\varepsilon_i \sim \text{Bernoulli}(1-p)$
 - ▶ First layer: $\mathbf{h} = \sigma((\mathbf{x} \odot \hat{\boldsymbol{\varepsilon}})\mathbf{W}_1)$
 - ▶ $(\mathbf{x} \text{diag}(\hat{\boldsymbol{\varepsilon}}))\mathbf{W}_1 = \mathbf{x}(\text{diag}(\hat{\boldsymbol{\varepsilon}})\mathbf{W}_1) = \mathbf{x}\hat{\mathbf{W}}_1$
 - ▶ Randomly setting to 0 rows of \mathbf{W}_1 (size $((D, L))$) with probability p



^adimension (1,D)

Dropout as a variational inference [Gal, 2016]

- **Illustration: dropout for a 2 layer NN** (1 hidden), $\epsilon_i \sim \text{Bernoulli}(1 - p_i)$:
 - ▶ $\mathbf{h}_1 = \sigma(\hat{\mathbf{x}}\mathbf{W}_1) = \sigma(\mathbf{x}\hat{\mathbf{W}}_1)$, $\hat{\mathbf{W}}_1 = \text{diag}(\hat{\epsilon}_1)\mathbf{W}_1$
 - ▶ $\hat{\mathbf{y}} =: f^{\hat{\mathbf{W}}_1, \hat{\mathbf{W}}_2}(\mathbf{x}) = \hat{\mathbf{h}}_1\mathbf{W}_2 = \mathbf{h}_1\hat{\mathbf{W}}_2$, $\hat{\mathbf{W}}_2 = \text{diag}(\hat{\epsilon}_2)\mathbf{W}_2$ - $\hat{\mathbf{W}} = \{\hat{\mathbf{W}}_1; \hat{\mathbf{W}}_2\}$
- **MC Dropout sampling:** $\frac{1}{S} \sum_{s \in S} p(\mathbf{y}_i | f^{\hat{\mathbf{W}}}(\mathbf{x}_i)) \approx \int p(\mathbf{y} | f^{\mathbf{W}}(\mathbf{x}^*)) q(\mathbf{W}) d\mathbf{w}$
 - ▶ \forall layer $l \in \{1; L\}$, \mathbf{W}_l random variable: $\mathbf{W}_l \sim q(\mathbf{W}_l) = g(\mathbf{M}_l, \epsilon_l) = \text{diag}(\epsilon_l)\mathbf{M}_l$
 - ▶ $\epsilon_{l,j} \sim \text{Bernoulli}(1 - p_l)$, \mathbf{M}_l deterministic parameters
 - ▶ $q(\mathbf{W}) = \prod_{l=1}^M q(\mathbf{W}_l)$
- **Big result** (see next): **training NN with dropout** \Leftrightarrow **training BNN with variational posterior approximation** $q_M(\mathbf{W})$ (and some prior $p(\mathbf{W})$)
 - ▶ $\mathbf{M} = \{\mathbf{M}_l\}_{l \in \{1; L\}}$ variational parameters
- **MC dropout:** sampling several passes with dropout \Leftrightarrow **performing MC approximate inference with variational posterior** $q_M(\mathbf{W})$

$$\frac{1}{S} \sum_{s \in S} p(\mathbf{y}_i | f^{\hat{\mathbf{W}}}(\mathbf{x}_i)) \approx \int p(\mathbf{y} | f^{\mathbf{W}}(\mathbf{x}^*)) q(\mathbf{W}) d\mathbf{w} \approx p(\mathbf{y} | \mathbf{x}^*, \mathbf{X}, \mathbf{Y})$$

Dropout as a variational inference [Gal, 2016]

Big result: proof sketch for a 2 layer NN (regression)

- **Prediction with dropout:** $\hat{\mathbf{y}} = \sigma(\mathbf{x}\hat{\mathbf{W}}_1)\hat{\mathbf{W}}_2 =: f^{\hat{\mathbf{W}}_1, \hat{\mathbf{W}}_2}(\mathbf{x})$
 - ▶ $\forall l \in \{1; 2\} : \hat{\mathbf{W}}_l = \text{diag}(\hat{\epsilon}_l)\mathbf{M}_l - \hat{\mathbf{W}} = \{\hat{\mathbf{W}}_1; \hat{\mathbf{W}}_2\}$
- **Training for regression,** $\hat{\mathcal{L}}_{dropout}$ objective function:

$$\hat{\mathcal{L}}_{dropout}(\mathbf{M}_1, \mathbf{M}_2) = \frac{1}{M} \sum_{i \in \mathcal{S}} \|f^{\hat{\mathbf{W}}}(\mathbf{x}_i) - y_i\|^2 + \lambda_1 \|\mathbf{M}_1\|^2 + \lambda_2 \|\mathbf{M}_2\|^2$$

- With Gaussian likelihood: $p(y_i | f^{\hat{\mathbf{W}}}(\mathbf{x}_i)) = \mathcal{N}(y, f^{\hat{\mathbf{W}}}(\mathbf{x}), \tau^{-1}\mathbf{I})$, we have:

$$\|f^{\hat{\mathbf{W}}}(\mathbf{x}_i) - y_i\|^2 = -\frac{1}{\tau} \log p(y_i | f^{g(\mathbf{M}, \hat{\epsilon}^i)}(\mathbf{x}))$$

- $\hat{\mathcal{L}}_{dropout}$ rewrites as follows:

$$\hat{\mathcal{L}}_{dropout}(\mathbf{M}_1, \mathbf{M}_2) = \frac{1}{M\tau} \sum_{i \in \mathcal{S}} \log p(y_i | f^{g(\mathbf{M}, \hat{\epsilon}^i)}(\mathbf{x})) + \lambda_1 \|\mathbf{M}_1\|^2 + \lambda_2 \|\mathbf{M}_2\|^2 \quad (1)$$

Dropout as a variational inference [Gal, 2016]

- Big similarity between $\hat{\mathcal{L}}_{dropout}$ in Eq (1) and algo 1!
- Same algorithms if:

$$\frac{\partial}{\partial M} KL(q(W)||p(W)) = \frac{\partial}{\partial M} N\tau(\lambda_1||M_1||^2 + \lambda_2||M_2||^2)$$

[Gal and Ghahramani, 2016] showed that this can be fulfilled for:

- $p(\mathbf{W}) = \prod_l p(\mathbf{W}_l) = \prod_l \mathcal{MN}(\mathbf{W}_l; 0; I/l_l^2, I)$ (prior factorized over layers)
- $q(\mathbf{W}_l) = \text{diag}(\hat{\epsilon}_l)M_l, \epsilon_{l,j} \sim \text{Bernoulli}(1 - p_j), q(\mathbf{W}) = \prod_l q(\mathbf{W}_l)$
 - ▶ Approximated by a mixture of two Gaussians with small std and one component fixed at zero
$$q_{\theta_{i,k}}(w_{i,k}) = (1 - p_i)\mathcal{N}(w_{i,k}; m_{i,k}; \sigma^2 I) + p_i\mathcal{N}(w_{i,k}; 0; \sigma^2 I)$$

⇒ A neural network with dropout can be interpreted as a variational Bayesian approximation

Model uncertainty

Predictive prediction with variational inference approximated with:

$$\begin{aligned} p(y|\mathbf{x}^*, \mathbf{X}, \mathbf{Y}) &= \int p(y|f^{\mathbf{W}}(\mathbf{x}^*))p(\mathbf{W}|\mathbf{X}, \mathbf{Y})d\mathbf{w} \\ &\approx \int p(y|f^{\mathbf{W}}(\mathbf{x}^*))q(\mathbf{W})d\mathbf{w} := q_{w^*}(y|\mathbf{x}^*) \end{aligned}$$

⇒ Estimate $p(y|\mathbf{x}^*, \mathbf{X}, \mathbf{Y})$ by MC sampling of $p(y|f^{\hat{\mathbf{W}}}(\mathbf{x}^*))$, $\hat{\mathbf{W}} \sim q(\mathbf{W})$

- $\mathbf{W} = \{W_l\}_{l=1}^L$ our set of random variables
- $f^{\mathbf{W}}(\mathbf{x}^*)$ our model's stochastic output
- $q_{W^*}(\mathbf{W})$ our optimum of variational distribution

Model uncertainty in regression

We will perform moment-matching and estimate the first two moments of the predictive distribution empirically.

Proposition

Given $p(y|f^{\mathcal{W}}(\mathbf{x}^*)) = \mathcal{N}(y; f^{\mathcal{W}}(\mathbf{x}^*); \tau^{-1}I)$ for some $\tau > 0$, $\mathbb{E}_{q_{\mathbf{w}^*}(y|\mathbf{x}^*)}[\mathbf{y}]$ can be estimated with the unbiased estimator

$$\tilde{\mathbb{E}}[\mathbf{y}] := \frac{1}{T} \sum_{t=1}^T f^{\hat{\mathbf{W}}_t}(\mathbf{x}^*) \xrightarrow{T \rightarrow \infty} \mathbb{E}_{q_{\mathbf{w}^*}(y|\mathbf{x}^*)}[\mathbf{y}]$$

with $\hat{\mathbf{W}}_t \sim q(\mathbf{W})$

⇒ equivalent to **performing T stochastic forward passes through the network and averaging the results.**

Model uncertainty in regression

Proposition

Given $p(y|f^W(\mathbf{x}^*)) = \mathcal{N}(y; f^W(\mathbf{x}^*); \tau^{-1}I)$ for some $\tau > 0$, $\mathbb{E}_{q_{w^*}(y|\mathbf{x}^*)}[(y)^T(y)]$ can be estimated with the unbiased estimator, with $\hat{W}_t \sim q(W)$:

$$\begin{aligned}\tilde{\mathbb{E}}[(y)^T(y)] &:= \tau^{-1}I + \frac{1}{T} \sum_{t=1}^T f^{\hat{W}_t}(\mathbf{x}^*)^T f^{\hat{W}_t}(\mathbf{x}^*) \\ &\xrightarrow{T \rightarrow \infty} \mathbb{E}_{q_{w^*}(y|\mathbf{x}^*)}[(y)^T(y)]\end{aligned}$$

Corollary

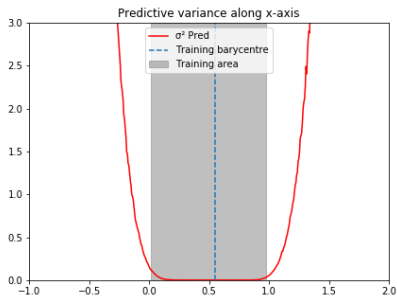
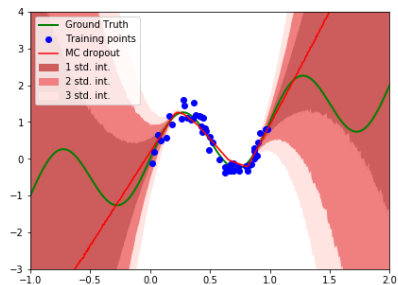
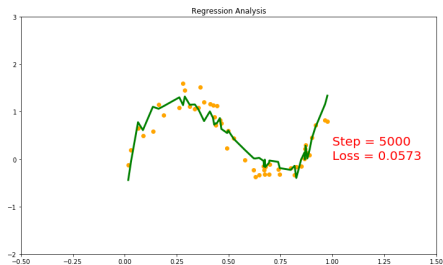
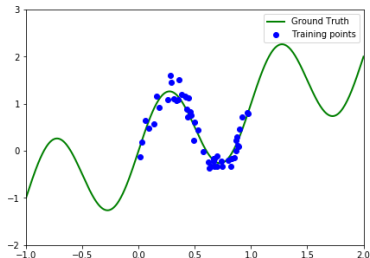
$\text{Var}_{q_{w^*}(y|\mathbf{x}^*)}[y]$ can be estimated with the unbiased estimator

$$\begin{aligned}\widetilde{\text{Var}}[(y)] &:= \tau^{-1}I + \frac{1}{T} \sum_{t=1}^T f^{\hat{W}_t}(\mathbf{x}^*)^T f^{\hat{W}_t}(\mathbf{x}^*) - \frac{1}{T} \left(\sum_{t=1}^T f^{\hat{W}_t} \right)^T \left(\sum_{t=1}^T f^{\hat{W}_t} \right) \\ &\xrightarrow{T \rightarrow \infty} \text{Var}_{q_{w^*}(y|\mathbf{x}^*)}[y]\end{aligned}$$

⇒ **sample variance of T stochastic forward passes** through the NN + the **inverse model precision**

Application: MC dropout for predictive distribution

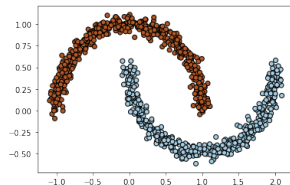
- MC dropout for regression: $\sin(x) + x$



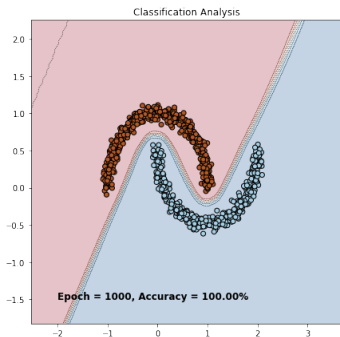
Application: MC dropout for predictive distribution

- MC dropout for non-linear classification

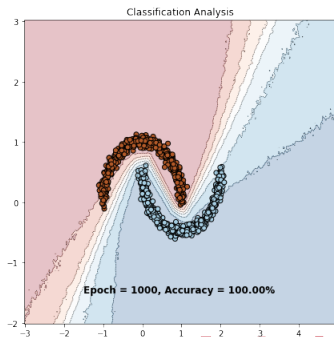
- As for BLR: $p(y = 1 | \mathbf{x}^*, \mathcal{D}) \approx \sum_{s=1}^S f_{\mathbf{w}^s}(\mathbf{x}^*)$, $\mathbf{w}^s \sim q(\mathbf{W})$



Deterministic NN



Bayesian NN (MC dropout)



References I

- [Bishop, 2006] Bishop, C. M. (2006).
Pattern Recognition and Machine Learning (Information Science and Statistics).
Springer-Verlag, Berlin, Heidelberg.
- [Blundell et al., 2015] Blundell, C., Cornebise, J., Kavukcuoglu, K., and Wierstra, D. (2015).
Weight uncertainty in neural network.
In Bach, F. and Blei, D., editors, *Proceedings of the 32nd International Conference on Machine Learning*,
volume 37 of *Proceedings of Machine Learning Research*, pages 1613–1622, Lille, France. PMLR.
- [Gal, 2016] Gal, Y. (2016).
Uncertainty in Deep Learning.
PhD thesis, University of Cambridge.
- [Gal and Ghahramani, 2016] Gal, Y. and Ghahramani, Z. (2016).
Dropout as a bayesian approximation: Representing model uncertainty in deep learning.
In *Proceedings of the 33rd International Conference on International Conference on Machine Learning - Volume 48, ICML'16*, pages 1050–1059. JMLR.org.
- [Graves, 2011] Graves, A. (2011).
Practical variational inference for neural networks.
In Shawe-Taylor, J., Zemel, R. S., Bartlett, P. L., Pereira, F., and Weinberger, K. Q., editors, *Advances in Neural Information Processing Systems 24*, pages 2348–2356. Curran Associates, Inc.
- [Hernandez-Lobato and Adams, 2015] Hernandez-Lobato, J. M. and Adams, R. (2015).
Probabilistic backpropagation for scalable learning of bayesian neural networks.
In Bach, F. and Blei, D., editors, *Proceedings of the 32nd International Conference on Machine Learning*,
volume 37 of *Proceedings of Machine Learning Research*, pages 1861–1869, Lille, France. PMLR.
- [Hinton et al., 2012] Hinton, G. E., Srivastava, N., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. (2012).
Improving neural networks by preventing co-adaptation of feature detectors.
CoRR, abs/1207.0580.
- [Hinton and van Camp, 1993] Hinton, G. E. and van Camp, D. (1993).
Keeping the neural networks simple by minimizing the description length of the weights.
In *Proceedings of the Sixth Annual Conference on Computational Learning Theory, COLT '93*, pages 5–13,
New York, NY, USA. ACM.

References II

- [Jylänki et al., 2014] Jylänki, P., Nummenmaa, A., and Vehtari, A. (2014).
Expectation propagation for neural networks with sparsity-promoting priors.
Journal of Machine Learning Research, 15:1849–1901.
- [Kingma and Welling, 2014] Kingma, D. P. and Welling, M. (2014).
Auto-encoding variational bayes.
In *2nd International Conference on Learning Representations, ICLR 2014, Banff, AB, Canada, April 14-16, 2014, Conference Track Proceedings*.
- [MacKay, 1992] MacKay, D. J. C. (1992).
A practical bayesian framework for backpropagation networks.
Neural Comput., 4(3):448–472.
- [Murphy, 2012] Murphy, K. P. (2012).
Machine Learning: A Probabilistic Perspective.
The MIT Press.
- [Neal, 1996] Neal, R. M. (1996).
Bayesian Learning for Neural Networks.
Springer-Verlag, Berlin, Heidelberg.